

COMP3310 Assignment 2

Alvin Tang (u7447218)

Due 26 April 2024

This PDF version is provided as a backup for the Markdown file (`README.md`).

1 Gopher Indexing Client

This Internet Gopher client written in the C programming language indexes directories and files. It is implemented in accordance with [RFC 1436](#). The overall design is to ensure that outbound requests strictly adhere to the protocol standard whilst handling inbound responses as gracefully as possible.

1.1 Initialisation

The program is intended for Linux machines with GCC installed. The standard C library is sufficient for compilation. Use the command `make` to compile the program with the source code. Simply execute the command `./client <hostname> <port>` to run the program. For example, if the Gopher server is hosted locally at port 70, `./client localhost 70` connects to it. An IP address can also be directly used. As an example, the command `./client 127.0.0.1 70` is also valid.

For local testing, a local Gopher server can be started using [Motsognir](#) with the command `sudo motsognir`. The listening port and the process can be listed using the command `sudo lsof -i :70`.

1.2 Program Design

The C programming language is used because of its extensive use in native Linux programs and its efficiency relative to alternatives (*e.g.*, Python, Java). A standard Linux installation should be able to execute the program; no additional packages are required.

The overall design relies on performing a breadth-first search (BFS) of the filesystem hosted on the Gopher server. The indexed files and directories are recorded in a linked list implemented using the `entry` struct. Each `entry` contains the type of the record (`item_type`), the record itself (`record`). The pointer `next` directs to the next entry in the linked list.

There are six types of records.

Entry type (<code>item_type</code>)	Record (<code>record</code>)
DIRECTORY	Pathname to the directory
TEXT	Pathname to the text file
BINARY	Pathname to the binary file
ERROR	Invalid request line
EXTERNAL	Request line referencing an external server
TIMEOUT	Pathname to the file taking excessive time to download
TOO_LARGE	Pathname to the file exceeding the size limit

Since the protocol is stateless and the connection is terminated once the server has responded, the function `gopher_connect()` establishes and closes a connection every time a new request is made. The function `gopher_connect()` appends “`\r\n`” to the request which is passed as the parameter `request`. It sends

the request line to the server. The function `func()`, passed as another argument to `gopher_connect()`, is responsible for receiving the response from the server.

Instantiations of `func()` include `indexing()`, `evaluate_file_size()` and `print_response()`. With reference to the protocol guidelines, the client listens for a response until the server terminates the connection. Therefore, these three functions use `recv()` provided by the Socket API in a while loop. The built-in `timeval` struct and the `errno` library in C are utilised for handling timeouts and failed responses. Notice that a failure to receive a response is handled differently from an empty string response from the server.

1.2.1 Recording Indexed Files and Directories

As long as the server adheres to the protocol's format, its response to a pathname is always the directory index with the first character of a row indicating the nature of the file. For instance, 0 indicates that the file is a (non-binary) text file whereas 1 refers to a subdirectory or an external server. The function `indexing()` uses this information to determine the type of the file/directory. A new entry is added to the linked list, including the item's pathname and type.

If the request is invalid, the corresponding response begins with 3. In this case, a new entry is added to the linked list if and only if it does not already exist. This enables us to count the number of invalid references at a later stage.

With reference to RFC 1436, the canonical type 9 refers to binary files. Actual server implementations often have more specific types and non-canonical types that gained popularity after RFC 1436 was published. In this client program, files that are not in plain text are considered binary.

Classification	File types
Directory	1
Reference to external server	1
Text file	0
Error	3
Binary file	4 (BinHex-encoded file), 5 (DOS file), 6 (uuencoded file) 9 (Binary file), g (GIF), I (Image), : (Bitmap) ; (Movie), < (Audio), d (Document), h (HTML) p (PNG), r (RTF), s (Sound), P (PDF), X (XML)

Any row in the response starting with the character `i` is a human-readable informational message, thus ignored by the indexation process. Other references such as Telnet (8 and T), CCSO nameserver (2) and mirror (+) do not fall within any of the categories above. These types are therefore disregarded by the indexation process.

1.2.2 Recursively Index Subdirectories

Following the indexation of the root directory using the request "`\r\n`", `main()` goes through the linked list and calls `gopher_connect()` for every subdirectory. This effectively models a breadth-first search of the filesystem hosted on the Gopher server.

1.2.3 Evaluation and Loading of File Content

Upon indexation of all directories and files in the filesystem, `evaluate()` is called to find the following information:

1. Number of directories, text files and binary files
2. Content of the smallest text file (if multiple exist, the first one indexed)
3. Sizes of the smallest and the largest text files
4. Sizes of the smallest and the largest binary files
5. Number of invalid references (each unique invalid request is counted once)

6. List of external server references (hostname and port) and their connectivity
7. List of references causing issues/errors

Note that if the server ends its response with “.\r\n” as per the protocol standard, the three characters are included in the size counts. This helps with determining whether the server is replying with proper responses.

The helper function `test_external_servers()` checks the connectivity to external servers by attempting to establish a connection. A server is considered “up” if it accepts the connection in a few seconds. Otherwise, it is marked as “down”.

1.2.4 Terminal Output

As the client program runs, logs are printed to the terminal (`stdout` and `stderr`). The terminal outputs to `stdout` are in the following formats.

- Request sent at <timestamp>: <request>
- Indexed <item type>: <pathname of newly indexed item>
- Number of <directories/text files/binary files/invalid references>: <number>
- Size of the <smallest/largest text/binary file>: <number>
- Other information produced by `evaluate()`

Issues and errors are printed to the terminal through `stderr`:

- Usage: <program name> <hostname> <port>
- Error: <error message>
- File too large: <pathname>
- Transmission timeout: <pathname>

1.2.5 Minimising Errors and Maximising Security

Whilst this client program follows the Gopher protocol standard in RFC 1436, measures are taken to reduce the impact of malformed responses from servers. For instance, the response is written to a string variable and is always explicitly terminated with the null terminator (“\0”). This avoids problems such as out-of-bound reading/writing.

Safer built-in functions in C (such as `strncpy()` rather than `strcpy()`) are used to avoid unexpected behaviour that might possibly be the result of long requests or responses. Indexed records in the linked list, with memory allocated using `malloc()`, are cleaned up by `cleanup()` before `main()` terminates.

At the beginning of `gopher_connect()`, the program is terminated if the connection cannot be established. An error message is printed to `stderr` specifying the issue.

1.2.6 Receiving Server’s Response through Multiple Packets

The built-in `recv()` function reads the responses from the server. The server might send it with multiple packets. In `indexing()`, for example, a while loop is used to read the response from the server and write it to the buffer until the server disconnects.

We set `BUFFER_SIZE` is 65536 bytes as this is a reasonable size limit for most directory indices on a Gopher server. It can be adjusted if the user of this program has specific needs.

1.2.7 Handling Empty Responses and Timeouts

A correctly implemented server should respond with some content under all circumstances. “Empty response from the server” is printed to the terminal if the server terminates the connection without any response. This clarifies to the user whether the connection is terminated or the client is still waiting for a response.

Servers may be busy or unresponsive at times. There are two types of timeouts.

1. The server accepts the connection but does not respond to a request.
2. The server accepts the connection and responds but it takes too much time.
3. The server does not accept the connection within the time limit.

Timeouts prevent the program from getting stuck indefinitely. The first situation is managed using the function `setsockopt()`, a built-in feature in the Socket API, and the `timeval` struct. The configuration is set by `gopher_connect()`. The limit is currently set as 10 seconds.

The second situation is handled using `select()` and `fcntl()`. The configuration is set by the function responsible for receiving the response. The limit is currently defined as 5 seconds.

The third situation applies in `test_external_servers()` with a 5-second limit.

1.2.8 Handling Edge Cases: Malformed or Non-Standard Responses

A well-implemented server’s directory index response should include four pieces of information, separated by tab characters. In case a line therein is malformed, the program handles it gracefully to avoid indexing invalid things. It is also designed to handle long responses from the server.

Motsognir, for example, ensures that every line including errors and informational messages includes three tab characters as separators.

```
3error fake fake 0
iInformational fake fake 0
```

Other servers format these responses differently. The class server, for instance, uses `invalid` rather than `fake`.

```
3error invalid 0
iInformation invalid 0
```

The client program ensures correct handling of responses when interacting with different servers.

A server’s response should end with a line containing “. \r\n” followed by the null terminator. However, if the server does not follow the standard protocol to the full extent and terminates transmission without “. \r\n”, the program still handles the response gracefully. The client program perceives a server’s termination of the connection as the end of data transmission. An example is `/misc/malformed2` on the class server.

Unicode was not published at the time the Gopher protocol was publicised. We assume that pathnames contain only ASCII characters, or else converted into ASCII representations on the server side (as it is in Motsognir). For example, a file named “goufør” is detected as “%C9%A1o%CA%8Af%C9%99r”.

Some malformed responses omit the hostname and the port in an entry of the directory index. Examples include `/misc/malformed1`. This client still attempts to index the file, assuming it exists on the server. If the file does not exist in the server, the pathname is regarded as an invalid reference.

1.2.9 Handling Edge Cases: Exceptionally Large Files

As part of the `evaluate()` function’s execution, the client has to download the files for analysis of their properties, such as their size. The files may sometimes be unexpectedly large. The global constant `FILE_LIMIT` specifies the maximum file size that the client may accept. This ensures that the client will not be flooded with infinite/unacceptable network traffic, of which malicious parties may take advantage. Files exceeding the size limit are not considered in the statistics (*e.g.*, sizes of largest files) evaluated by `evaluate()`.

Currently, `FILE_LIMIT` is set as 2^{17} (131,076) bytes, considering that most files on `comp3310.ddns.net` do not exceed this limit. In realistic situations, users of this client should know the size scale of files that the target server hosts. This global constant can be adjusted accordingly. Files exceeding the file limit are logged and printed before the client program execution ends.

1.3 Testing

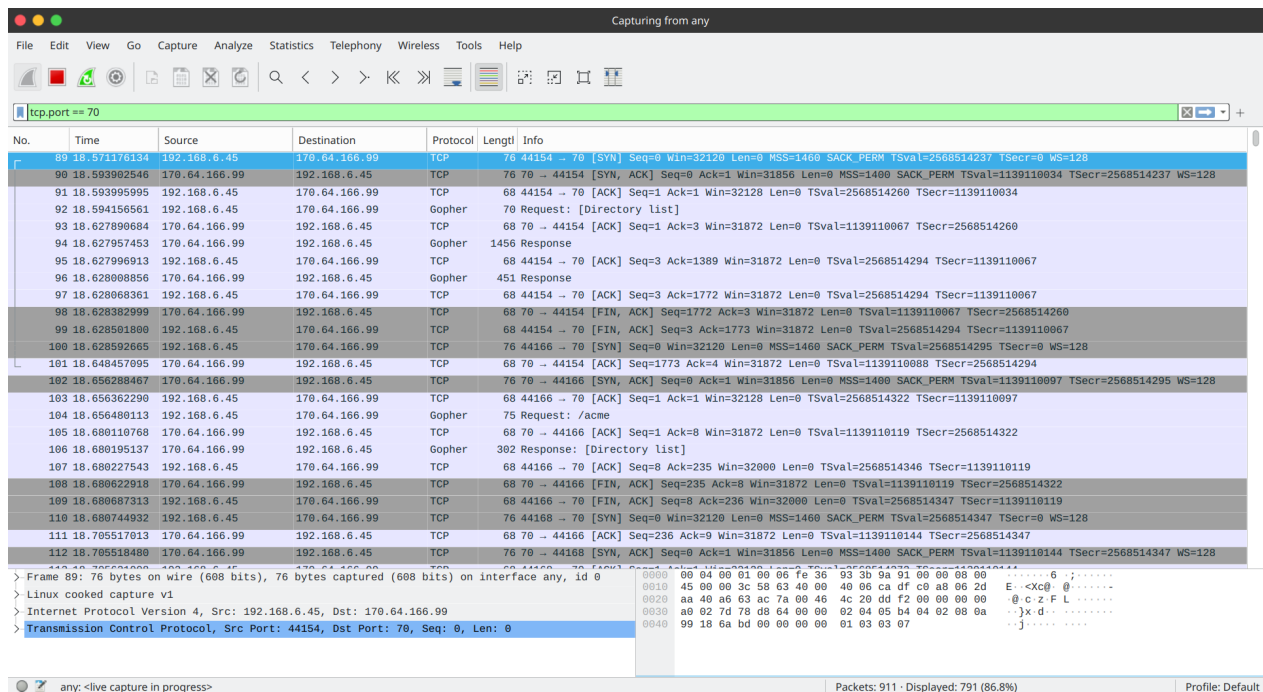
The client program is tested with the class server and a local Motsognir server. The terminal outputs are included for reference:

- `assets/output1.txt` — Class server, 2^{17} -byte file size limit.
- `assets/output2.txt` — Local Motsognir server, no file size limit.

The program is also tested with `gopher.floodgap.com` at port 70. The first 1024 lines of terminal output are included in `assets/output3.txt` for information.

Wireshark is used for monitoring the network traffic. The initial request is detected as “Request: [Directory List]”. There are occasional retransmissions of TCP packets and spurious acknowledgement (ACK) signals.

In the screenshot below, the first three rows show the establishment of the TCP connection. The fourth row is the Gopher request for the root directory index. The sixth and eighth rows correspond to the response from the server, delivered with two packets. For each TCP packet received, the client sends an ACK signal to the server.



References

- [1] F. Anklesaria, M. McCahill, P. Lindner, D. Johnson, D. Torrey, and B. Alberti. RFC 1436: The Internet Gopher Protocol, March 1993. URL <https://www.rfc-editor.org/rfc/rfc1436>.
- [2] caf. C: socket connection timeout. Stack Overflow, April 2010. URL <https://stackoverflow.com/a/2597774>.
- [3] dgookin. The `gettimeofday()` Function. C For Dummies Blog, July 2020. URL <https://c-for-dummies.com/blog/?p=4236>.
- [4] dhanshreekulkarni21. `strcat()` in c. GeeksforGeeks, March 2023. URL <https://www.geeksforgeeks.org/strcat-in-c>.

- [5] Foobar. Client in C++, use `gethostbyname` or `getaddrinfo`. Stack Overflow, October 2018. URL <https://stackoverflow.com/q/52727565>.
- [6] Frédéric Hamidi. What does **EAGAIN** mean? Stack Overflow, October 2010. URL <https://stackoverflow.com/a/4058377>.
- [7] IEEE and The Open Group. `setsockopt`, 2003. URL <https://linux.die.net/man/3/setsockopt>.
- [8] Arkaitz Jimenez. Time stamp in the C programming language. Stack Overflow, September 2009. URL <https://stackoverflow.com/a/1444440>.
- [9] Milan. How to convert string to IP address and vice versa. Stack Overflow, March 2011. URL <https://stackoverflow.com/a/5328184>.
- [10] Basile Starynkevitch. How to use `select()` on sockets properly? Stack Overflow, September 2015. URL <https://stackoverflow.com/a/32711642>.
- [11] Tomek Szpakowicz. When to use static keyword before global variables? Stack Overflow, December 2009. URL <https://stackoverflow.com/a/1856642>.
- [12] Wikipedia contributors. Gopher (protocol). Wikipedia. URL [https://en.wikipedia.org/wiki/Gopher_\(protocol\)](https://en.wikipedia.org/wiki/Gopher_(protocol)).